Three ways to run the same team

A short explainer for review before anything gets built. 29 October 2026.

The job

You say: "*Make me an ad for Apple.*"
The team gives back a **shot list**: what happens in each shot, and where the camera and the light go.

No pictures and no video are made. It stops at the written shots.

The team (the same in all three)

- **The Interrogator** (AI)
  Reads your idea and the website, and asks about what's missing. You can always say "you decide".
- **The Cinematographer** (AI)
  Picks the camera and the light for every shot.
- **The Director** (AI)
  Writes the story: the place, the person, the product, and each shot.
- **The Ruler and the Glue** (plain code, not AI)
  The Ruler checks the rules. The Glue sticks everything into the final prompt.

The one question we are testing

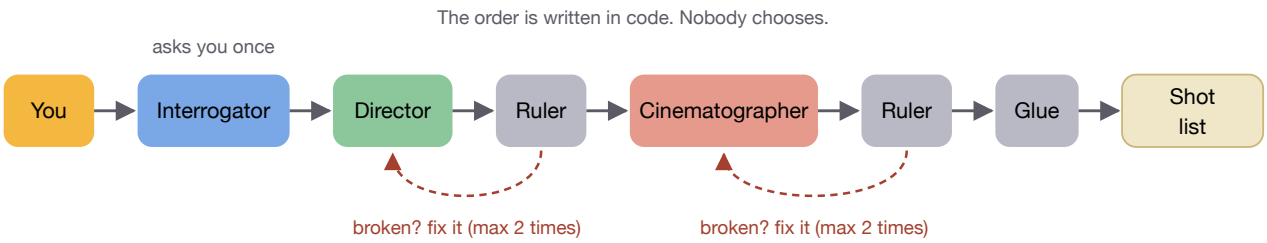
Who decides who works next?

- **Sequential:** a fixed recipe written in code. Like a factory line.
- **Coordinator:** one boss AI tells the helpers what to do. Like a teacher with helpers.
- **Swarm:** the helpers pass the work to each other. Like passing a ball.

The same helpers, the same instructions, the same Ruler, and the same "two tries to fix it" rule in all three. Only the wiring changes, so any difference comes from the wiring.

1. Sequential: the factory line

what we have today

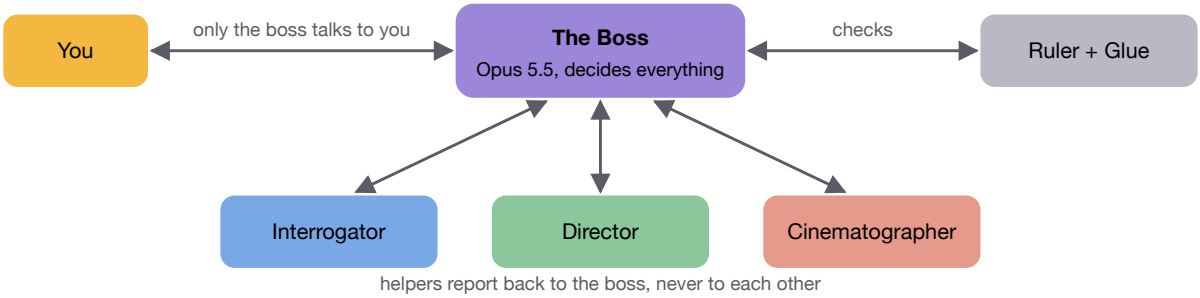


Good: always the same steps. Cheap, fast and easy to debug.

Bad: it can't change plans. If a brief needs something new, the recipe doesn't know.

Boss: the code. The helpers run on Sonnet 5.5 at high effort.

2. Coordinator: the teacher with helpers

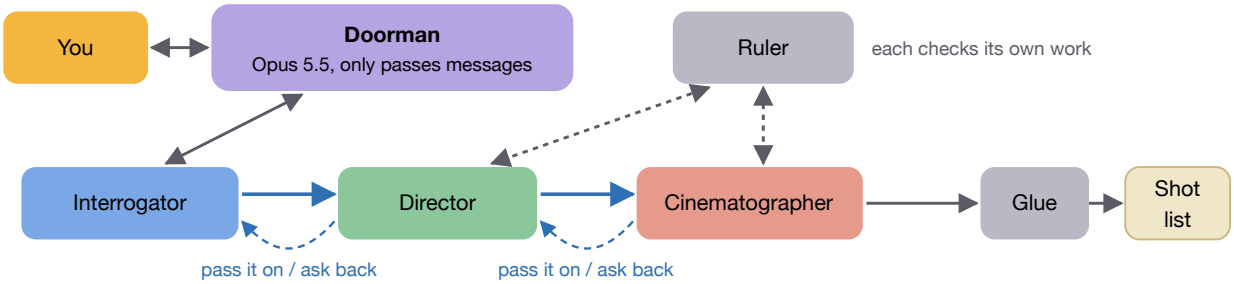


Good: flexible. The boss can skip, repeat or change the order when a brief is strange.

Bad: costs more and is slower. A helper only knows what the boss tells it, and the boss can choose badly.

Boss: Opus 5.5 at high effort. The helpers run on Sonnet 5.5 at high effort.

3. Swarm: passing the ball



Good: helpers can ask each other directly, like the Cinematographer asking the Director where the window is.

Bad: nobody is really in charge, so the work can drift or loop. It's also the hardest to follow and costs the most.

Claude Code always keeps one lead, so ours is a Doorman (Opus 5.5, high effort) who only passes your messages and never gives orders. The helpers run on Sonnet 5.5 at high effort. Best watched in iTerm2, one pane per helper.

Side by side

	Sequential	Coordinator	Swarm
Who decides	The code	The boss AI	Each helper
Helpers talk to	Nobody (they hand over results)	Only the boss	Each other
Where you watch it	Claude Code, step by step	Claude Code, one row per helper	iTerm2, one pane per helper
My guess	Most reliable, cheapest	Worth it only for strange briefs	Most to learn from, most likely to drift

How you'll use it

- Type **/direct** in Claude Code. Pick sequential, coordinator or swarm. Give your idea and, if you have one, a website.
- Answer the Interrogator's questions once, or say "you decide". Then watch.
- You get the shot list, plus numbers: time, tokens, how many fixes were needed, and whether it passed.
- We run the same 3 ideas through all 3 teams (9 runs), then write the real report for your teacher.

A run only counts as a success if every check passed. A run that stopped is never shown as a win.